

Serverless means ...

AWS
Lambda

Serverless means ...



**No servers to provision
or manage**



Scales with usage



Pay for value



**Availability and fault
tolerance built in**

Serverless means:

Serverless means:

Greater agility

Less overhead

Better focus

Increased scale

More flexibility

Faster time to market

Today's focus:



Event-driven compute

Functions as a service

Serverless FaaS



Lambda Handles

Load Balancing

Auto Scaling

Handling Failures

Security Isolation

OS Management

Managing Utilization

(and many other things) for you


AWS Lambda release history



*As of October 2018, does not include region launches

Serverless applications

Event source



Function



Services (anything)



Serverless applications

Event source




Changes in
data state



Requests to
endpoints



Changes in
resource state



Function



Services (anything)



Serverless applications

Event source



Changes in
data state



Requests to
endpoints



Changes in
resource state



Function



Node.js
Python
Java
C#
Go
Ruby
Runtime API

NEW!

Services (anything)



Anatomy of a Lambda function

Handler() function

Function to be executed upon invocation

Event object

Data sent during Lambda function Invocation

Context object

Methods available to interact with runtime information (request ID, log group, more)

```
public String handleRequest(Book book, Context context) {  
    saveBook(book);  
  
    return book.getName() + " saved!";  
}
```

Introducing: AWS Lambda runtime API and layers



Features that allow developers to share, discover, and deploy both libraries and languages as part of their serverless applications



Runtime API enables developers to use Lambda with any programming language.



Layers let functions easily share code. Upload layer once, reference within any function.




Layers promote separation of responsibilities, lets developers focus on writing business logic.



Combined, runtime API and layers allow developers to share any programming language or language version with others

Fine-grained pricing



Free Tier

1M requests and 400,000 GBs of compute.
Every month, every customer.

Buy compute time in 100ms increments

Low request charge

No hourly, daily, or monthly minimums

No per-device fees

Never pay for idle

Tweak your function's computer power



Lambda exposes only a memory control, with the **% of CPU core and network capacity** allocated to a function proportionally

Is your code CPU, Network or memory-bound? If so, it could be **cheaper** to choose more memory.

Smart resource allocation

Match resource allocation (up to 3 GB!) to logic

Stats for Lambda function that calculates 1000 times all prime numbers
<= 1000000

| | | |
|---------|--------------|------------|
| 128 MB | 11.722965sec | \$0.024628 |
| 256 MB | 6.678945sec | \$0.028035 |
| 512 MB | 3.194954sec | \$0.026830 |
| 1024 MB | 1.465984sec | \$0.024638 |

Green==Best

Red==Worst

Smart resource allocation

Match resource allocation (up to **3 GB!**) to logic

Stats for Lambda function that calculates **1000 times** all prime numbers
<= 1000000

128 MB

256 MB

512 MB

1024 MB

-10.256981sec

+\$0.000001

Green==Best

Red==Worst

av

Lambda execution model

Synchronous (push)



Asynchronous (event)



Poll-based



Lambda API



API provided by the Lambda service

Used by all other services that
invoke Lambda across all models

Supports sync and async

Can pass any event payload
structure you want

Client included in every SDK


Lambda permissions model

Fine grained security controls for both execution and invocation:

- **Execution policies:**
- Define what AWS resources/API calls can this function access via IAM
- Used in streaming invocations
- E.g. "Lambda function A can read from DynamoDB table users"
- **Function policies:**
- Used for sync and async invocations
- E.g. "Actions on bucket X can invoke Lambda function Z"
- Resource policies allow for cross account

```
1 {  
2   "Version": "2012-10-17",  
3   "Statement": [  
4     {  
5       "Effect": "Allow",  
6       "Action": [  
7         "logs:CreateLogGroup",  
8         "logs:CreateLogStream",  
9         "logs:PutLogEvents"  
10      ],  
11      "Resource": "*"   
12    }  
13  ]  
14 }
```

AWS Account A



AWS Account B



Common Lambda use cases



Web Applications

- Static websites
- Complex web apps
- Packages for Flask and Express

Backends

- Apps & services
- Mobile
- IoT

Data Processing

- Real time
- MapReduce
- Batch

Chatbots

- Powering chatbot logic


Amazon Alexa

- Powering voice-enabled apps
- Alexa Skills Kit

IT Automation

- Policy engine
- Extending AWS services
- Infrastructure management

Amazon API Gateway



Create a unified
API frontend for
multiple micro-
services



DDoS protection
and throttling for
your backend

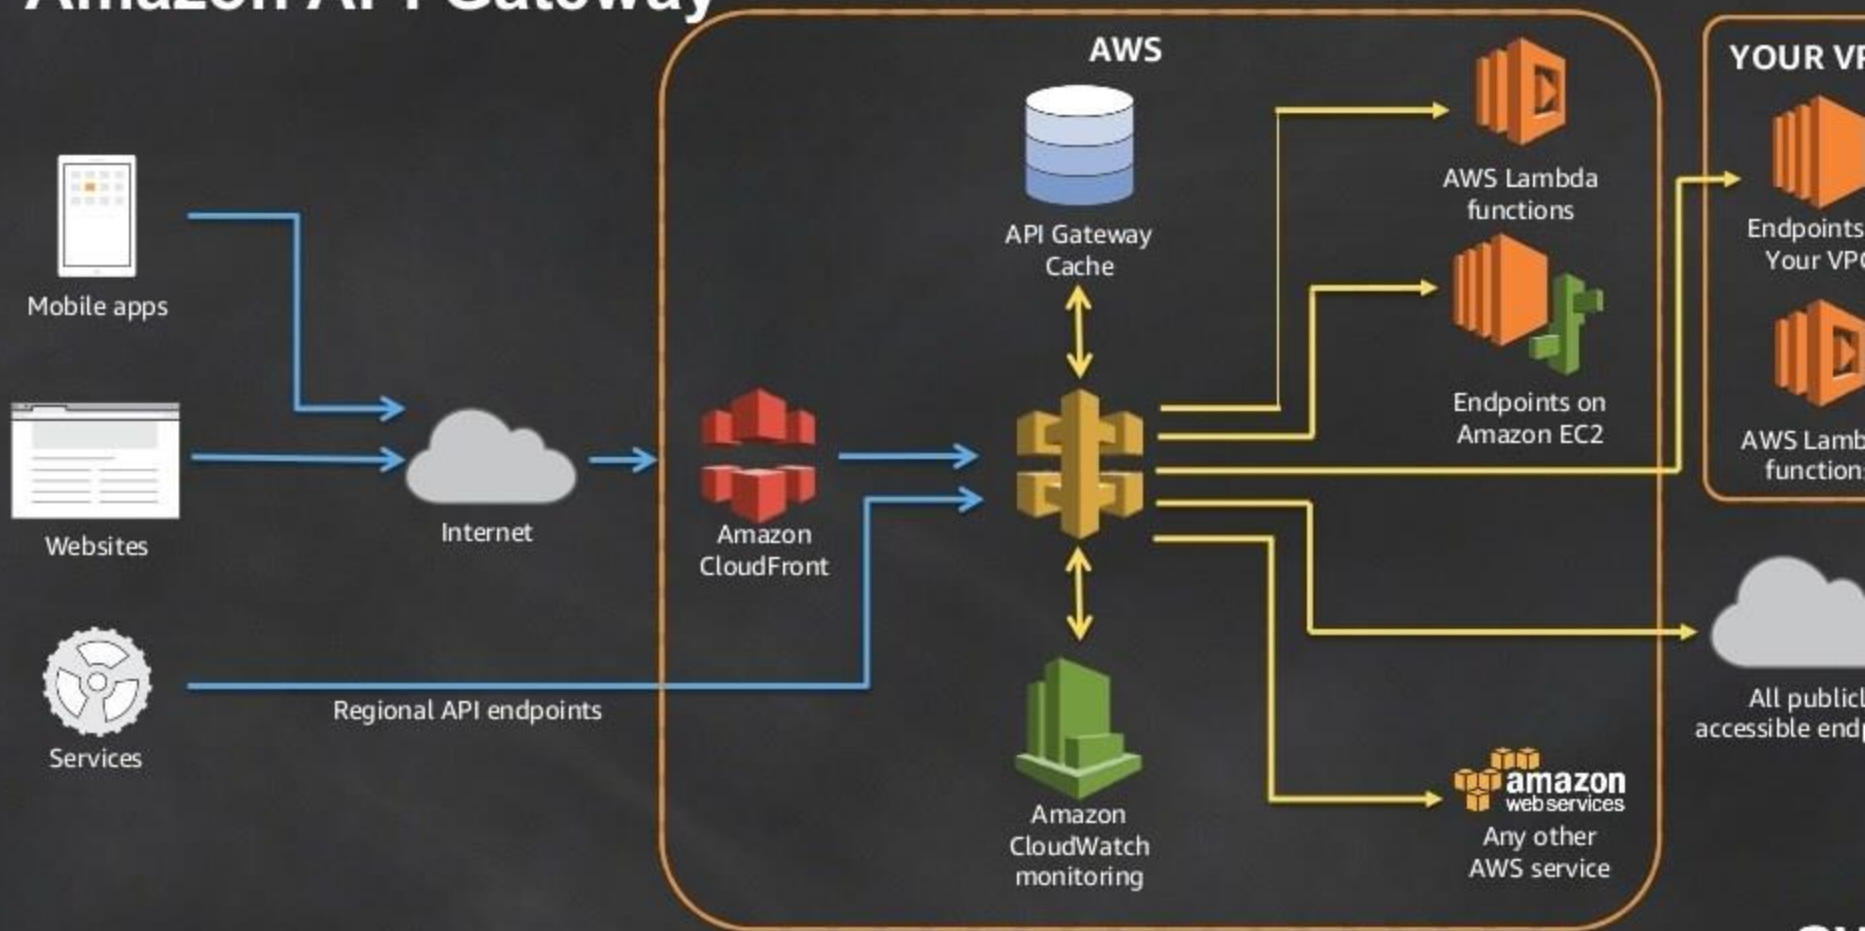


Authenticate and
authorize
requests to a
backend




Throttle, meter,
and monetize API
usage by third-
party developers

Amazon API Gateway



Introducing: API Gateway WebSockets

Coming soon!




Build real-time two-way communication applications chat, alerts and notifications, and streaming dashboards

Fully managed APIs to handle connections and messages transfer between users and backend services

Invoke AWS services like Lambda, Kinesis, or any HTTP endpoint based on message content

Pay for what you use based on connection management and messages transferred

Serverless web application with API Gateway

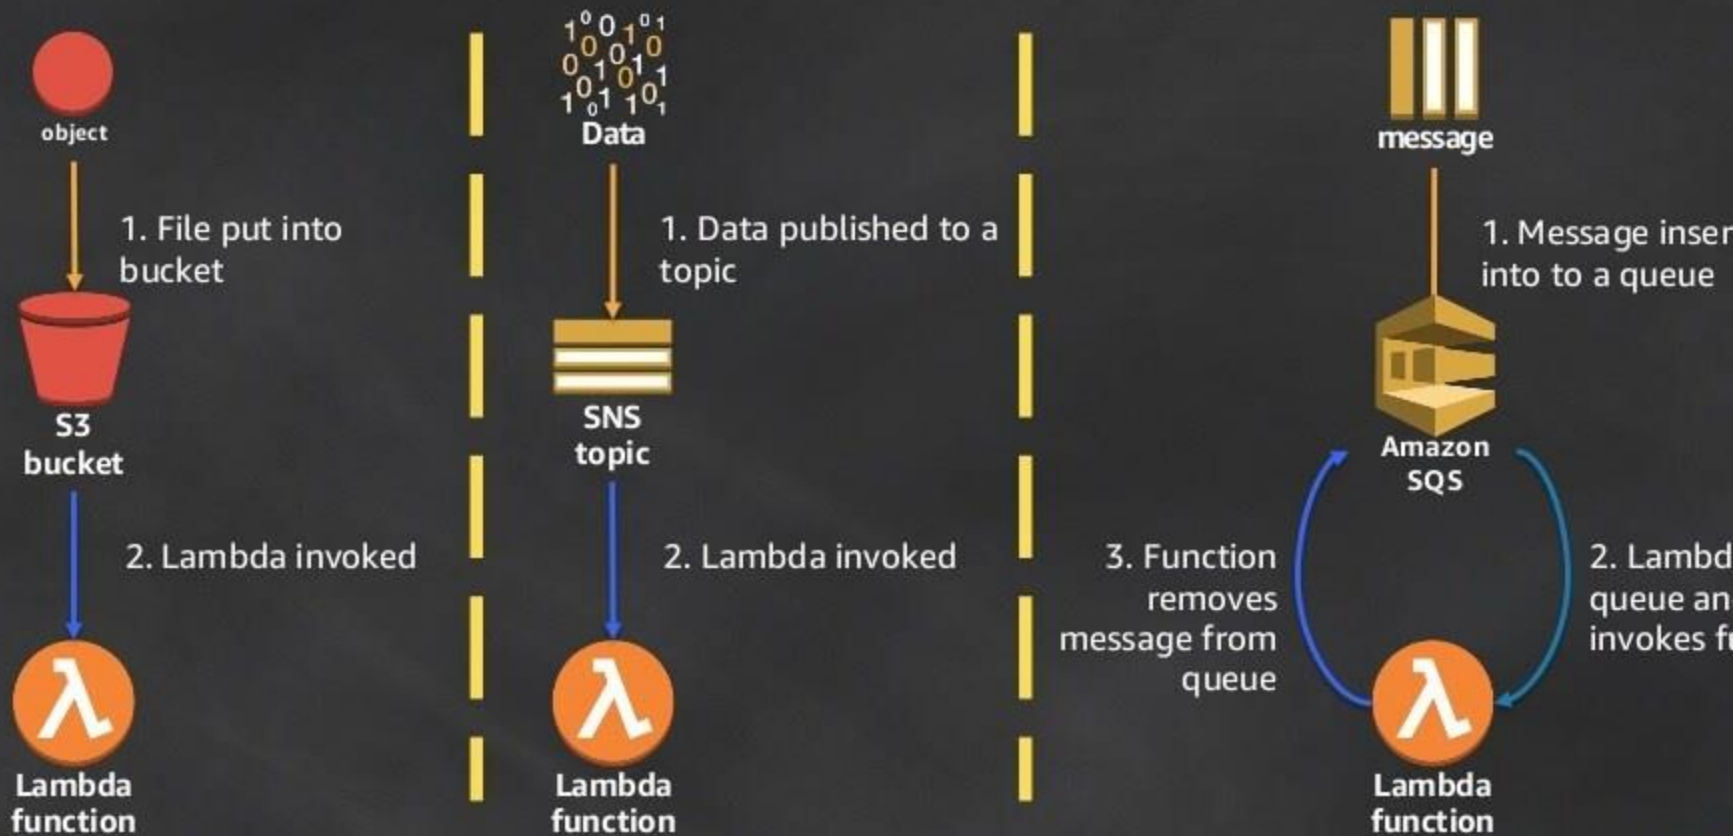


Amazon Simple Storage Service (Amazon S3) stores all of your static content: CSS, JS, images, and more. You would typically front this with a CDN such as CloudFront.


API Gateway handles all your application routing. It can handle authentication and authorization, throttling, DDOS protection, and more.

Lambda runs all the logic behind your website and interfaces with databases, other backend services, or anything else your site needs.

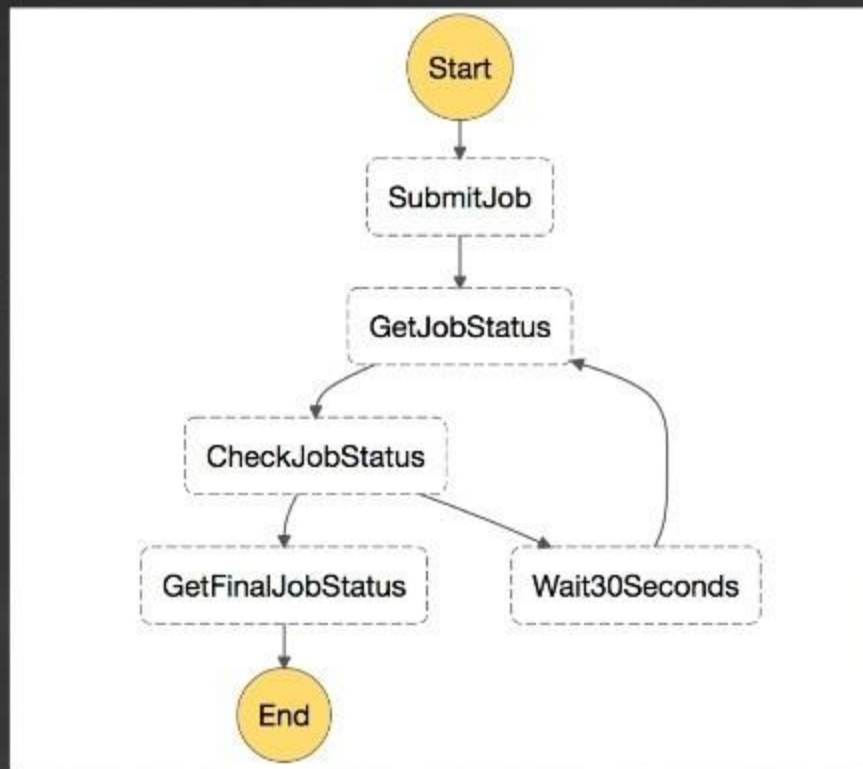
Serverless architectures



Serverless architectures



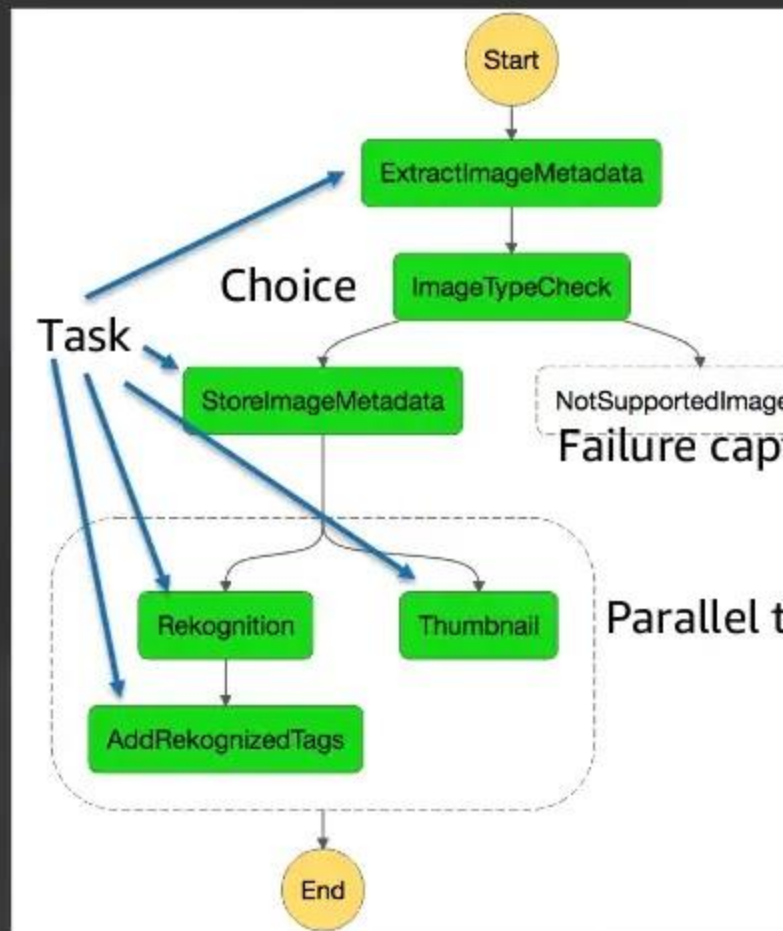
Keep orchestration out of code



AWS Step Functions

“Serverless” workflow management with zero administration

- Makes it easy to coordinate the components of distributed applications and microservices using visual workflows
- Automatically triggers and tracks each step and retries when there are errors, so your application executes in order and as expected
- Logs the state of each step, so when things do go wrong, you can diagnose and debug problems quickly



Build PCI and HIPAA compliant serverless applications!



Serverless platform services that can be used in both:



AWS
Lambda



Amazon
S3



Amazon
CloudFront



Amazon
DynamoDB



Amazon
Kinesis
Data
Streams



Amazon
Cognito



Amazon API
Gateway



Amazon
SNS

Metrics and logging are a universal right!

- **CloudWatch Metrics:**
- **7 Built in metrics for Lambda**
 - Invocation Count, Invocation duration, Invocation errors, Throttled Invocation, Iterator Age, DLQ Errors, Concurrency
 - Can call “**put-metric-data**” from your function code for custom metrics
- **7 Built in metrics for API-Gateway**
 - API Calls Count, Latency, 4XXs, 5XXs, Integration Latency, Cache Hit Count, Cache Miss Count
 - Error and Cache metrics support *averages and percentiles*



Metrics and logging are a universal right!

- **CloudWatch Logs:**
- **API Gateway Logging**
 - 2 Levels of logging, ERROR and INFO
 - Optionally log method request/body content
 - Set globally in stage, or override per method
- **Lambda Logging**
 - Logging directly from your code with your language's equivalent of `console.log()`
 - Basic request information included
- **Log Pivots**
 - Build metrics based on log filters
 - Jump to logs that generated metrics
- **Export logs to AWS ElastiCache or S3**



Messaging Systems in AWS

Messaging Enables Decoupling


Message Queueing

- Asynchronous
- Point-to-point



Publish-subscribe (pub-sub)

- Broadcast
- Point-to-multipoint



AWS Messaging Services

Amazon
SQS



- Simple, fast, reliable, managed message queues
- Build highly scalable decoupled applications

Amazon
SNS



- Simple, fast, reliable, managed pub/sub
- Highly scalable push messaging to users or applications

Amazon
Kinesis



- Highly scalable streaming data analysis
- Ingest and analyze real-time data

AWS
IoT



- Managed service to connect devices to the cloud
- Supports billions of subscribers

Amazon
Pinpoint



- Mobile notification and engagement platform
- Targeted campaigns over email, SMS, and mobile

Amazon
MQ



- Managed message broker service for Apache ActiveMQ

Amazon SQS: key features



- Persistent message queue with high **durability** and **availability**



- Messages are stored across **multiple AZs**
- Messages retained until deleted— **up to 14 days**



- Nearly unlimited **throughput**

Amazon SQS: key features



- Amazon CloudWatch metrics and alerts



- Message payloads up to 256 KB (5 TB using Amazon S3)



- Message batching to increase throughput and reduce cost

- Secure: uses AWS Identity & Access Management (IAM) and HTTPS/TLS



Two SQS Queue Types

Standard Queues

- At least once delivery
- Best-effort ordering
- Nearly unlimited transaction rate

FIFO Queues

- First-in First-out delivery preserving message
- Exactly once processing
- 300 transactions per second
- Messages groups for



SQS Extended Client Library

- For sending message payloads larger than 256KB
- Message payloads are stored in an S3 bucket
- SQS is used to transmit a reference to the S3 payload
- Abstracted from developer with Java client-side library

Amazon SNS: key features



- Proven **reliability** with messages are stored across **multiple AZ**



- **Flexible** message delivery over multiple transport protocols



- Nearly unlimited **throughput**

Amazon SNS: key features



- **Instantaneous** or delayed, push-based delivery



- **Simple** APIs and easy integration



- **Amazon CloudWatch** metrics and alerts



- Message payloads up to **256 KB**

Amazon MQ: Key Features



JMS

AMQP

STOMP



NMS

MQTT

WebSocket

Compatible with industry-standard APIs and protocols.
Managed message broker service for Apache ActiveMQ.

Amazon MQ: Key Features

- Transient & persistent messaging
- Local & distributed transactions (XA)
- Queues & topics (with FIFO)
- Composite & virtual destinations
- Message filtering
- Request/reply
- Scheduled messages
- Unlimited message size
- Unlimited message retention



FAQ – when should I use Amazon MQ vs SQS/SNS?

SQS & SNS

- For **cloud-native** applications
- Simple
- Unlimited throughput
- Fully managed

Amazon MQ Service

- **For application migration**
- **API-compatible**
- **Feature-rich**
- **Limited scale**
- **Managed infrastructure**