

Amazon Elastic Map Reduce: the concepts

A horizontal blue banner with a white background for the text. The banner is decorated with various white line-art icons including clouds, lightning bolts, cubes, gears, and lightbulbs, scattered across its surface.

What to expect

- Amazon EMR
- Amazon S3 as HDFS
- Core node and Task nodes
- Elastic clusters
- Beyond Map Reduce: Spark
- Q&A

What Hadoop is good at

Semi-structured/unstructured data

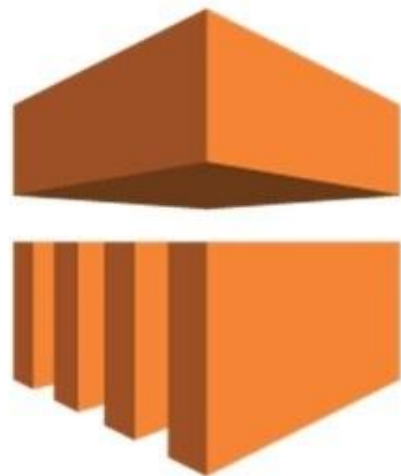
Disparate Data Sets

ETL at scale

Batch Analytics

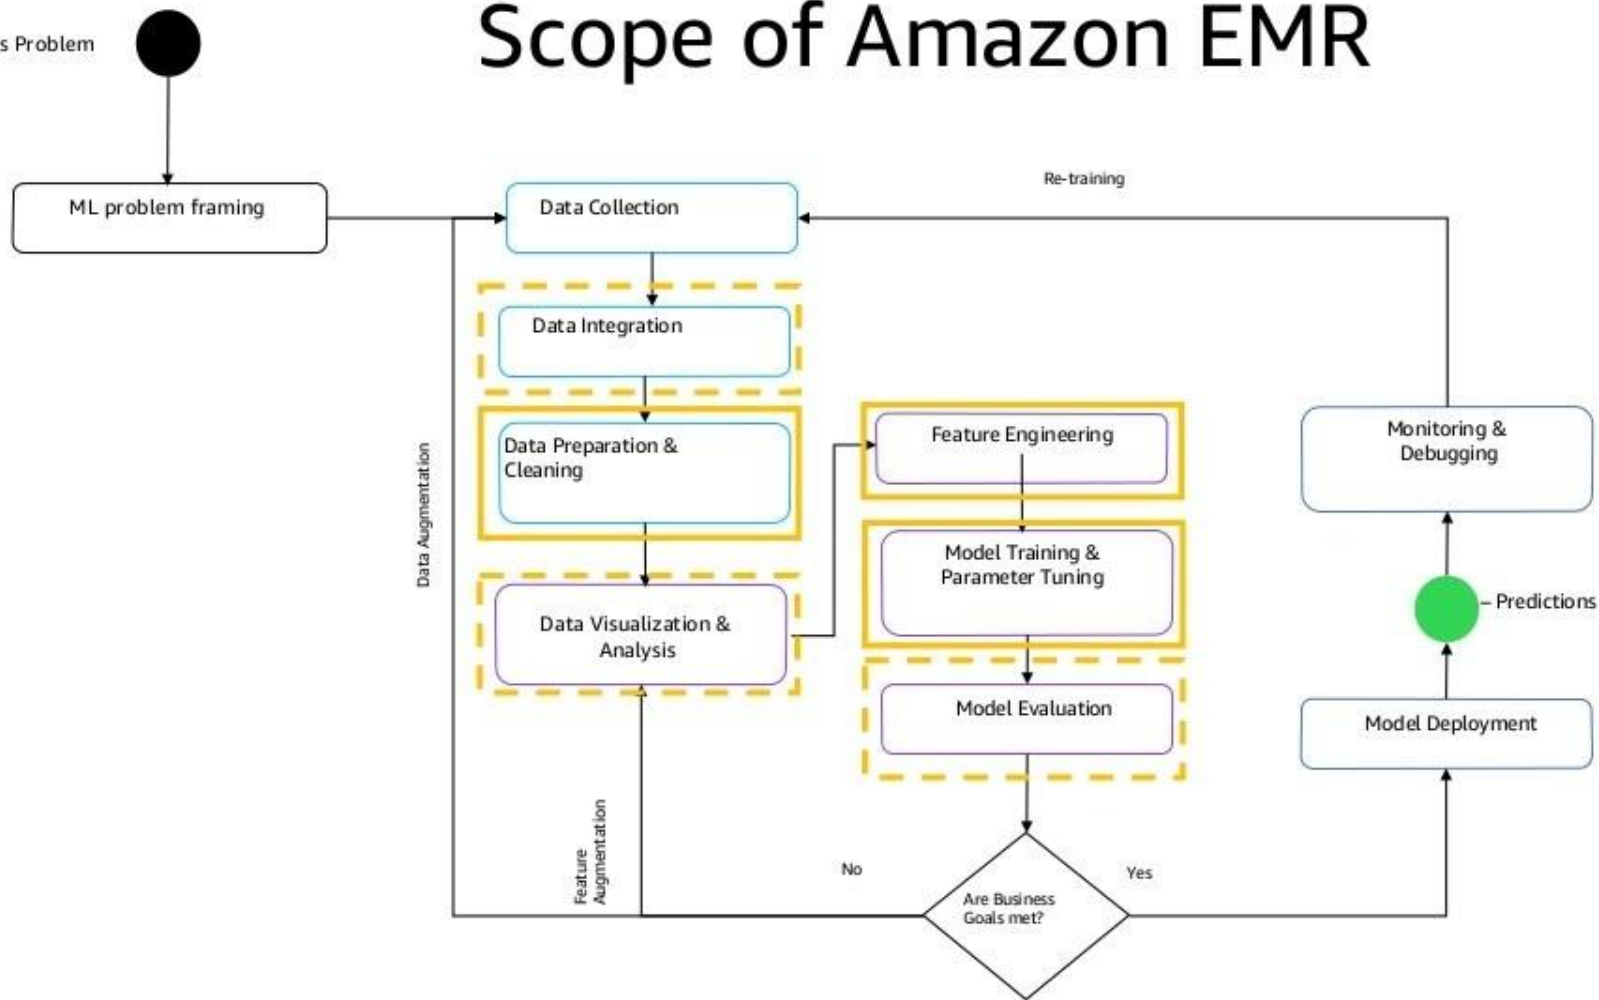
Log Processing & Aggregation

Amazon EMR – Hadoop, Spark, Presto in the Cloud



- Managed platform
- Launch a cluster in minutes
- Leverage the elasticity of the cloud
- Baked in security features
- Pay by the hour and save with Spot
- Flexibility to customize

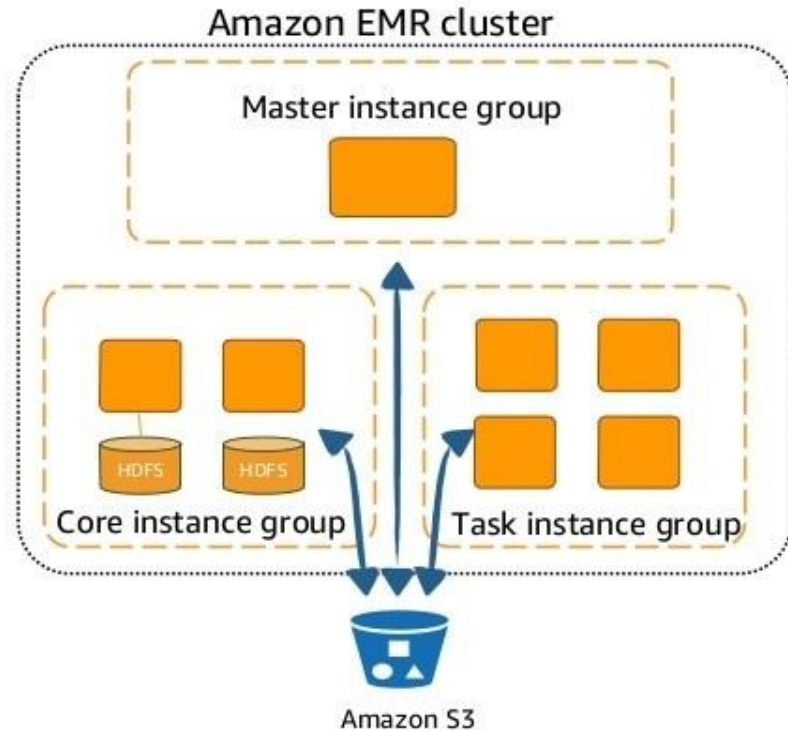
Business Problem



Amazon S3 as HDFS

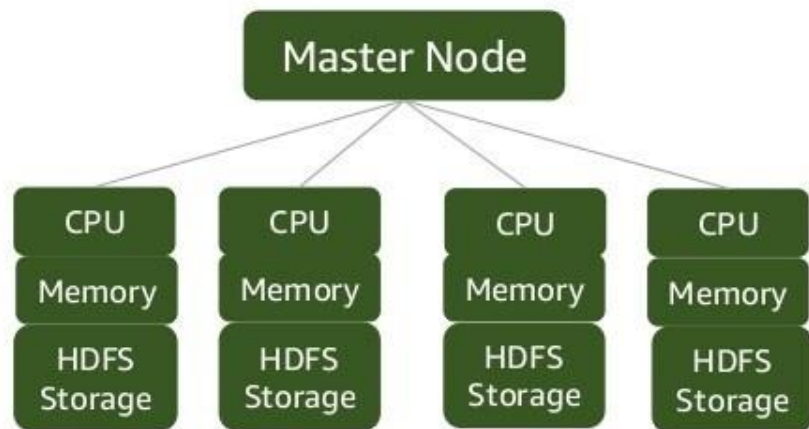
Amazon S3 as HDFS

- Use Amazon S3 as your **permanent** data store
- HDFS for **temporary** storage data between jobs
- No additional step to copy data to HDFS



Understanding Decoupled Storage & Compute

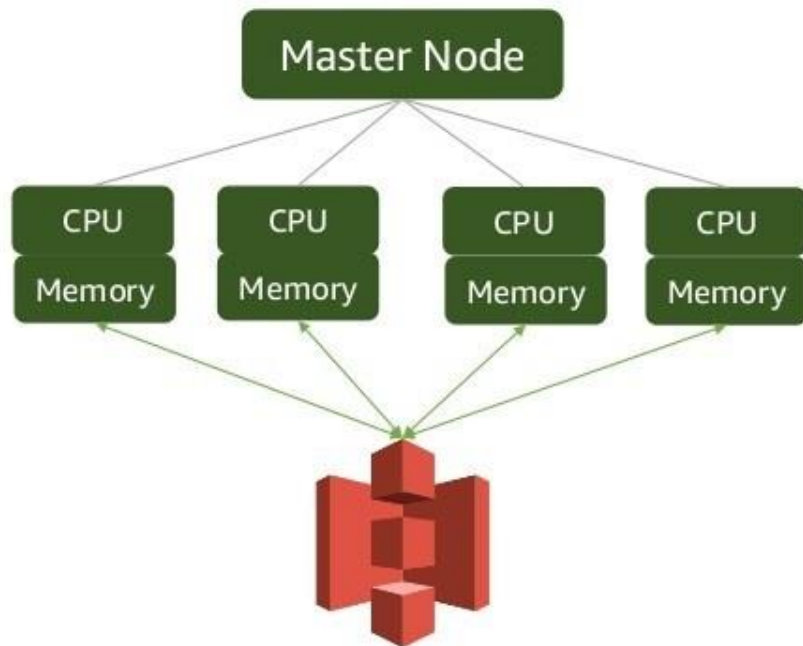
Old Clustering / Localized Model



HDFS = 3X Replication

500 TB Dataset = 1.5 PB cluster with replication

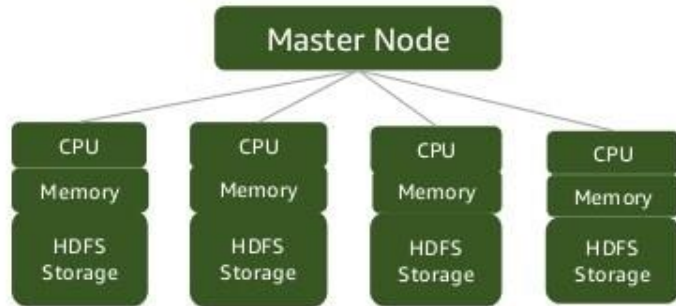
Amazon EMR Decoupled Model



S3 as Streaming HDFS via EMRFS

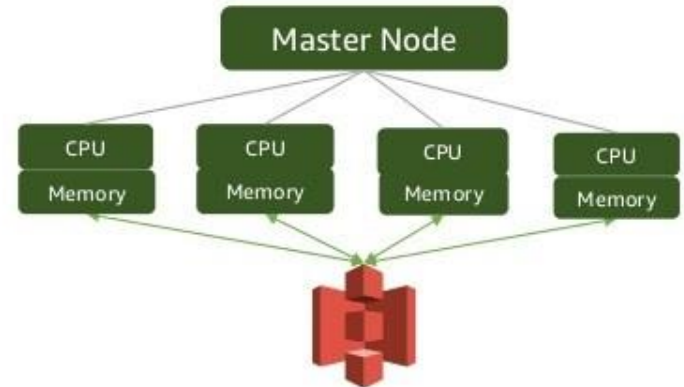
Understanding Decoupled Storage & Compute

Hadoop on EC2



- 3X HDFS Replication of Data across Nodes
- Protected against data loss from node failure
- Single AZ replication
- Paying for storage units via EC2 ephemeral drives or EBS volumes
- Data stored locally so cluster = 24x7 by default

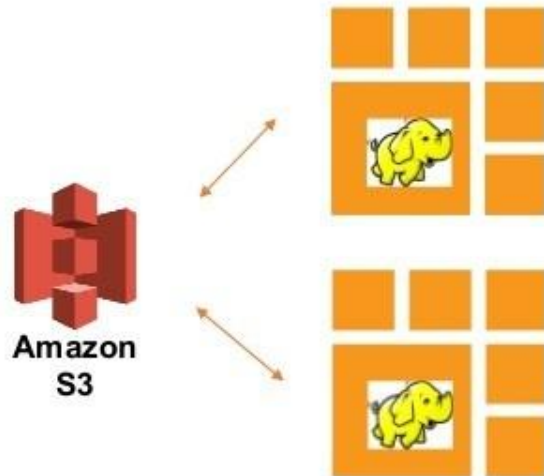
EMR



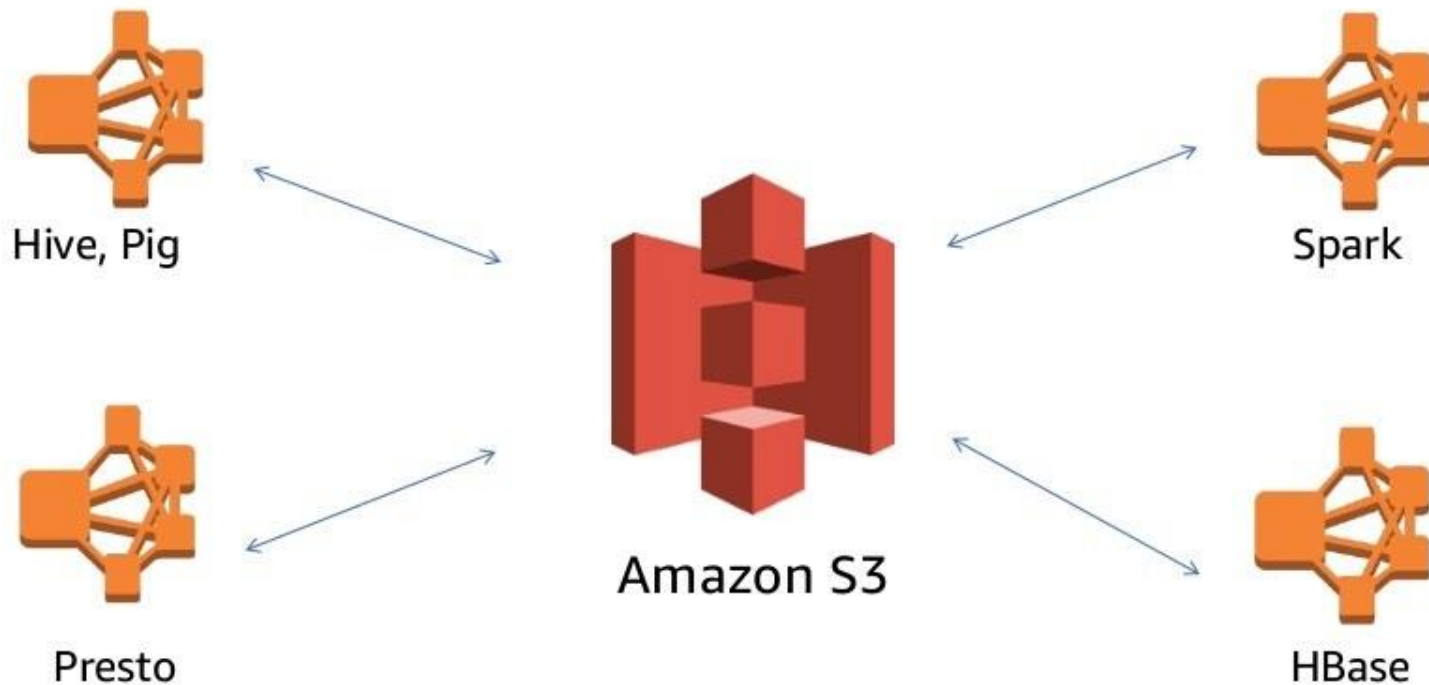
- Native S3 multi-AZ replication
- Protected against data loss from node failure, cluster failure or AZ failure
- Multiple physical facility replication
- Paying for storage units via S3 (cheap!!!)
- Spin clusters up and down or turn off!
- Stream multiple EMR clusters from same dataset in S3

Amazon S3 as your persistent data store

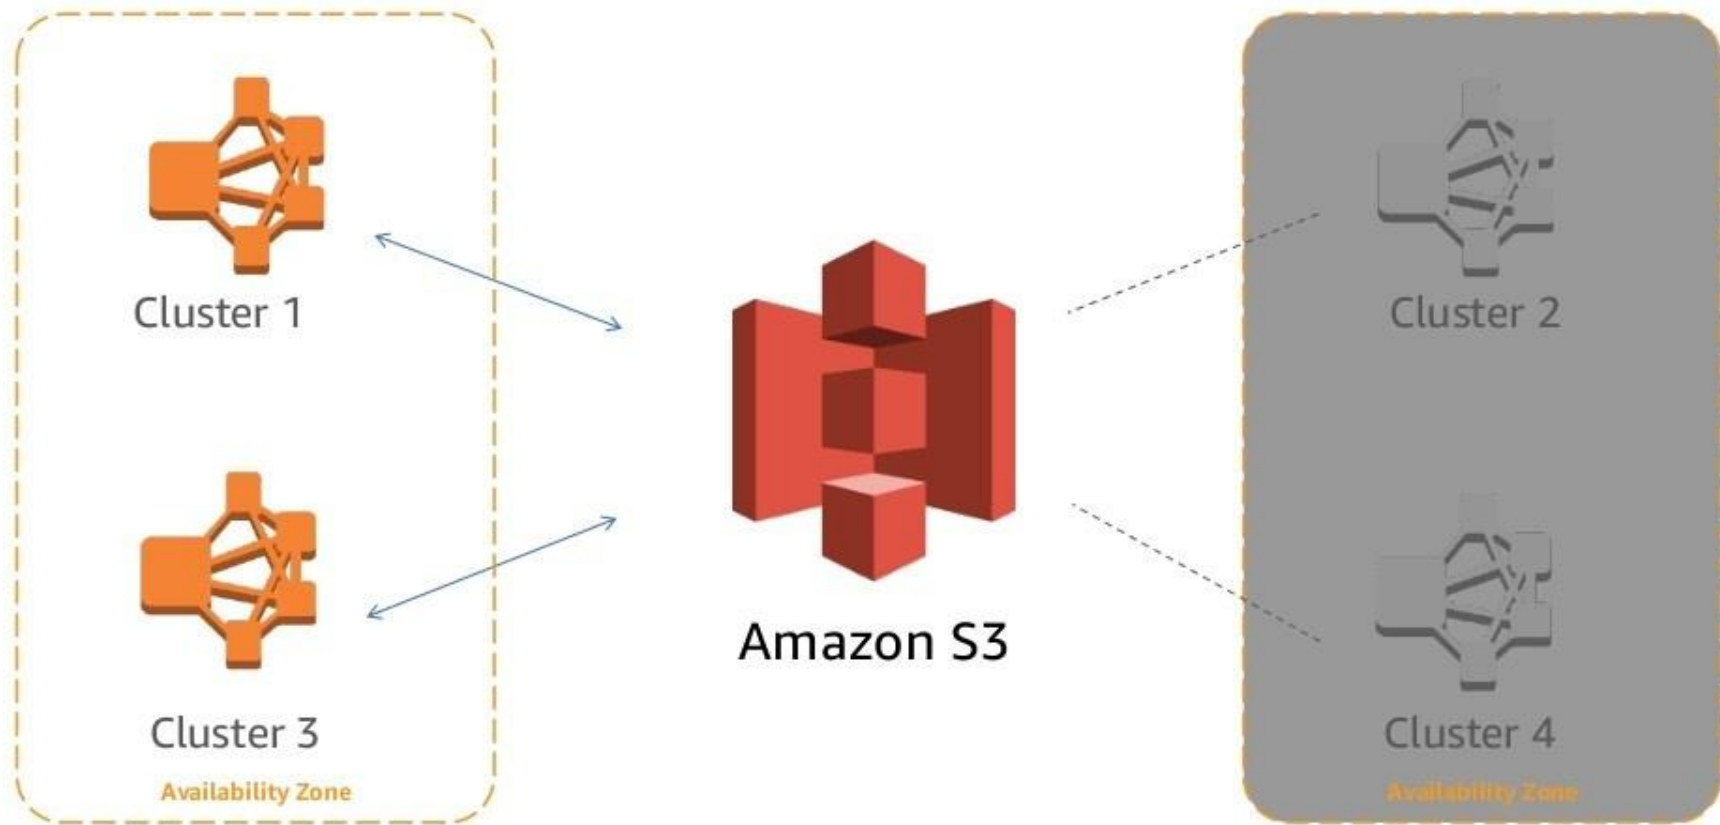
- Designed for 99.999999999% durability
- Separate compute and storage
- Resize and shut down clusters with no data loss
- Point multiple clusters at the same data in S3
- Easily evolve your analytics infrastructure as technology evolves



EMR with Amazon S3

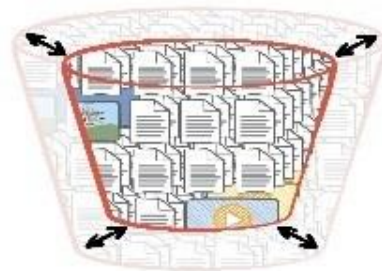


Disaster Recovery built in



Benefits: Amazon S3 as HDFS

- No need to scale HDFS
 - Capacity?
 - Replication?
- Amazon S3 scales with your data
 - Both in **IOPS** and **storage**
 - Your data can grow **independent** of CPU and RAM



Benefits: Amazon S3 as HDFS

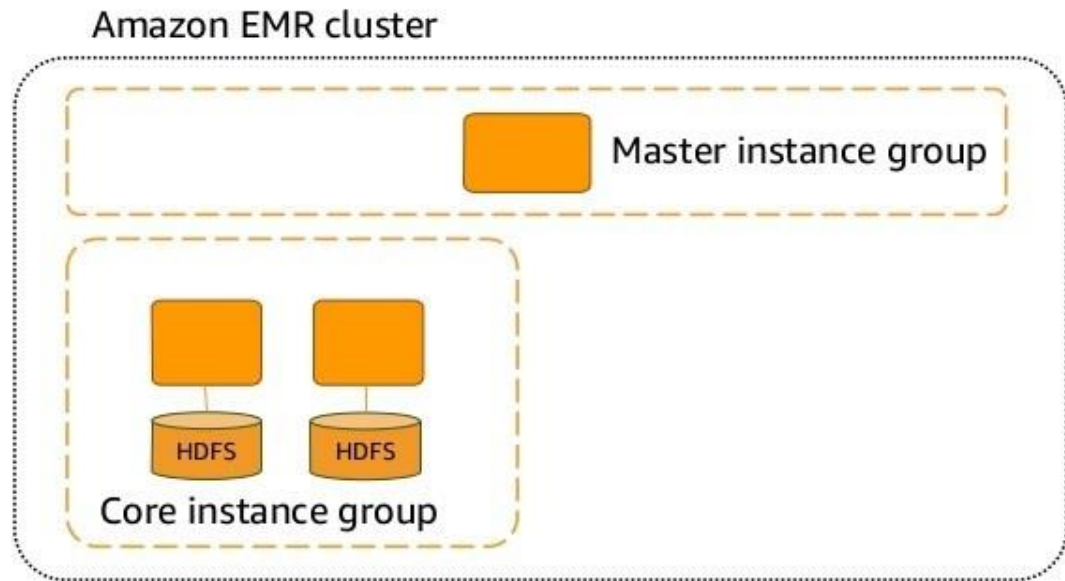
- Take advantage of **S3 features**
 - Server-side encryption
 - Lifecycle policies
 - Versioning to protect against corruption
- Build **elastic** clusters
 - Add nodes to read from Amazon S3
 - Remove nodes with data safe on Amazon S3

Core nodes and task nodes

Amazon EMR Core nodes

Run TaskTrackers
(Compute)

Run DataNode
(HDFS)

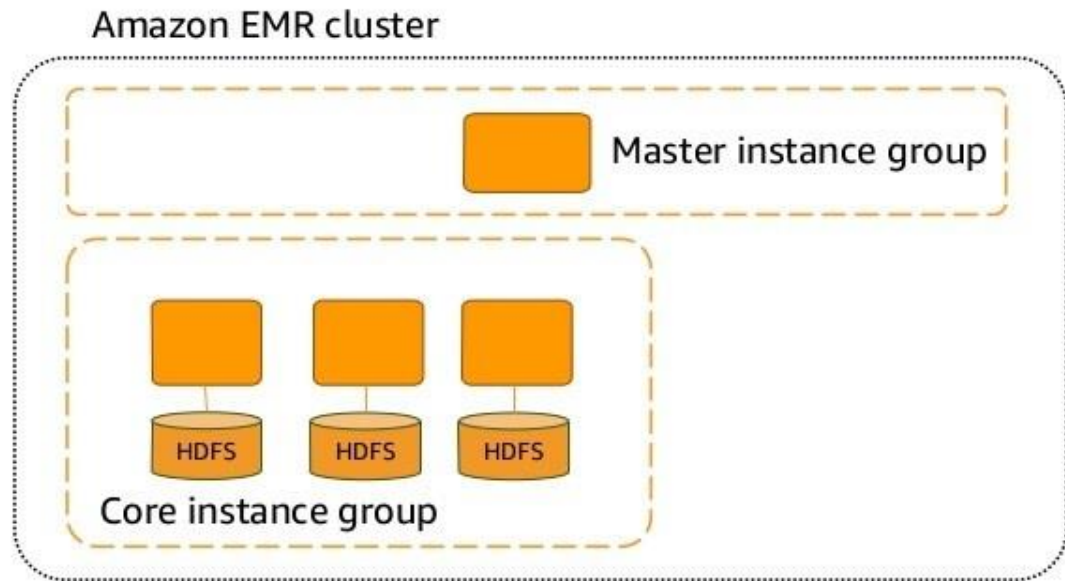


Amazon EMR Core nodes

Can add core nodes

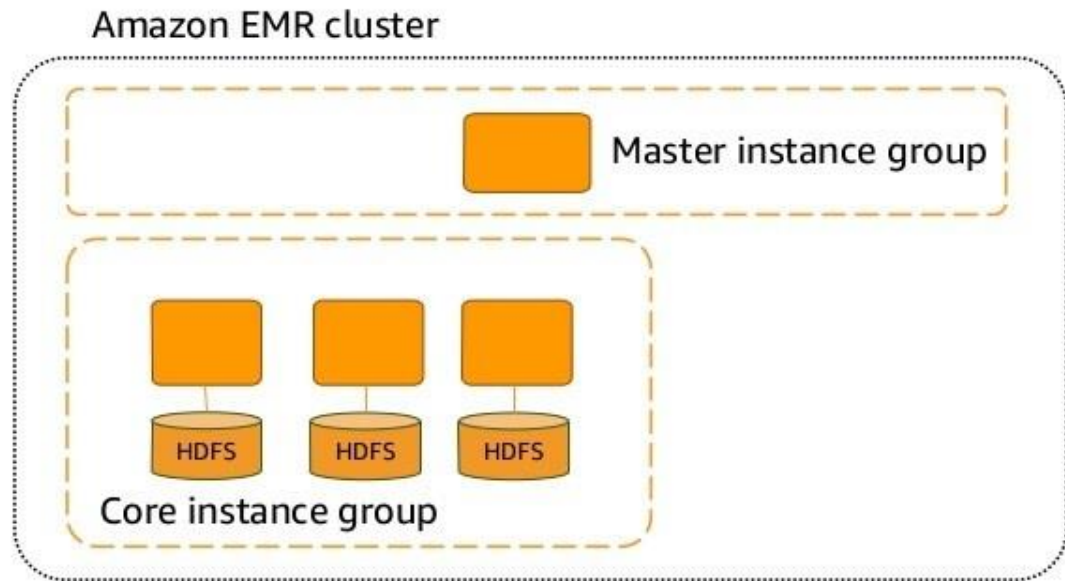
More HDFS space

More CPU/memory



Amazon EMR Core nodes

Can't remove core nodes because of HDFS

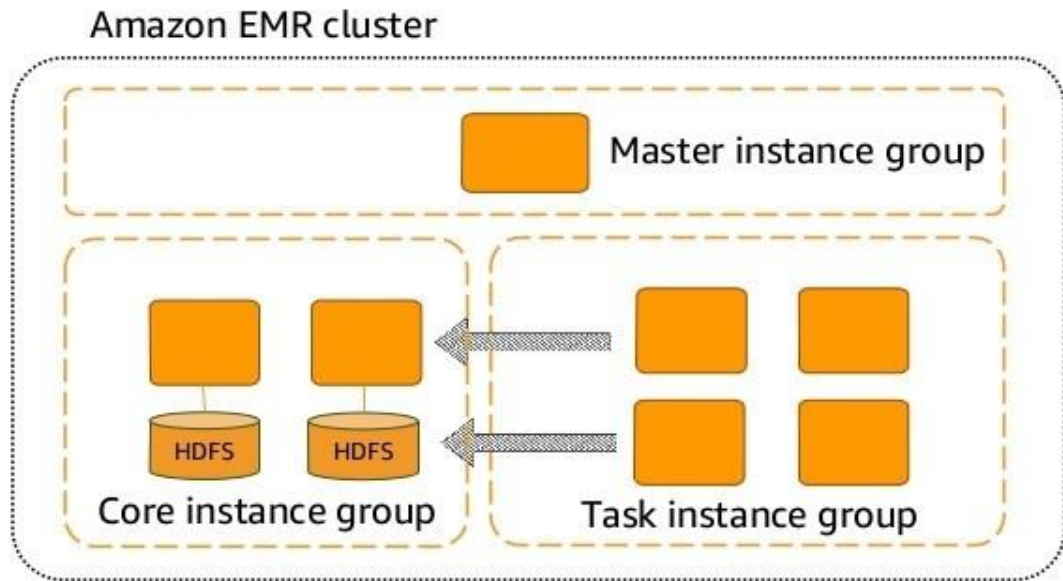


Amazon EMR Task nodes

Run TaskTrackers

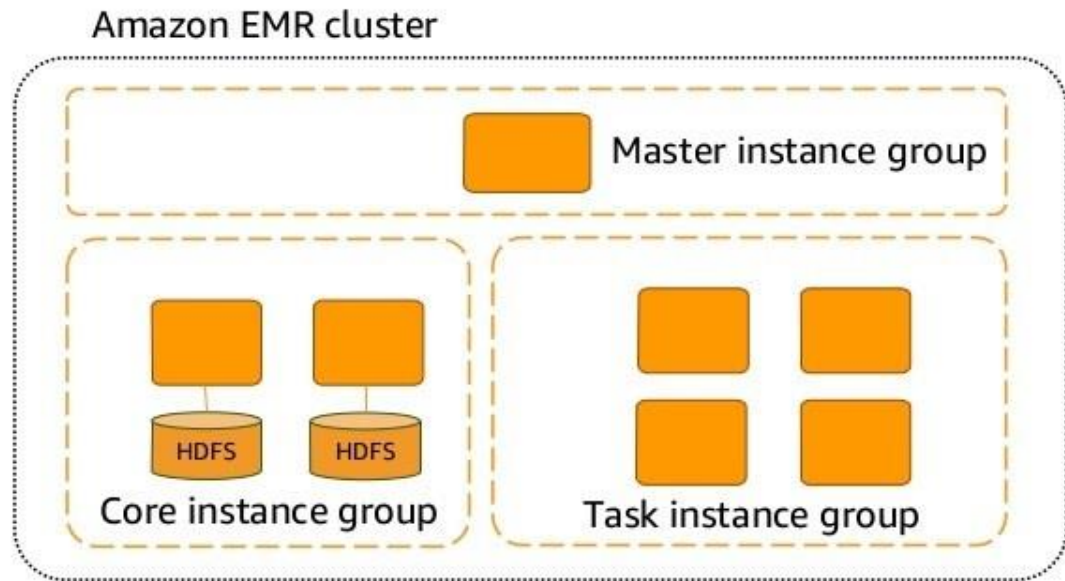
No HDFS

Reads from core nodes



Amazon EMR Task nodes

Can add task nodes

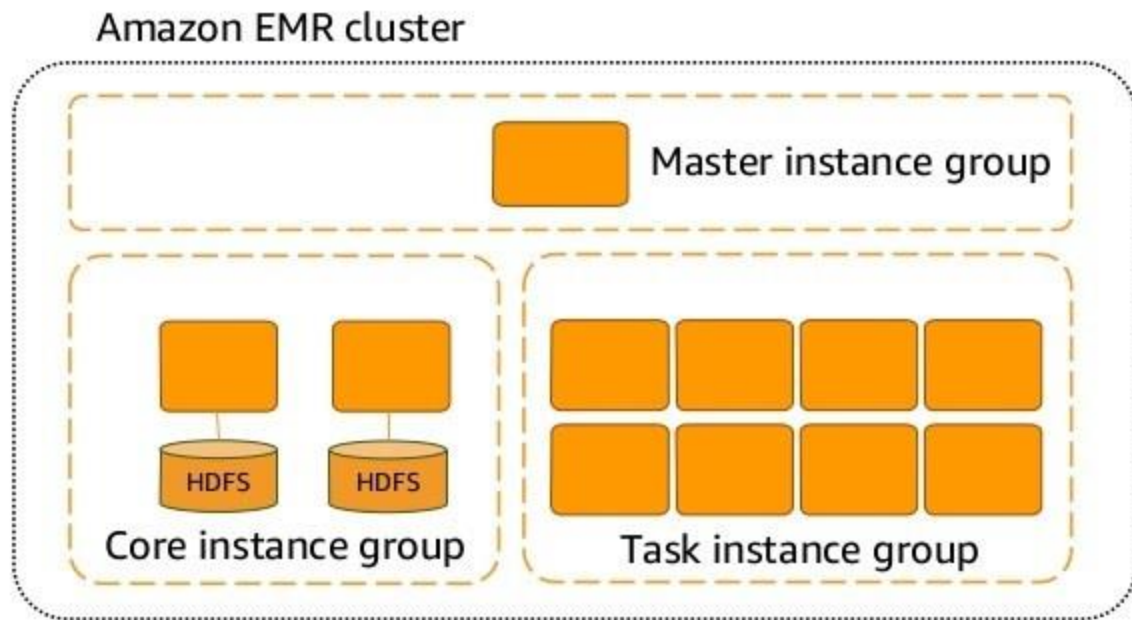


Amazon EMR Task nodes

More CPU power

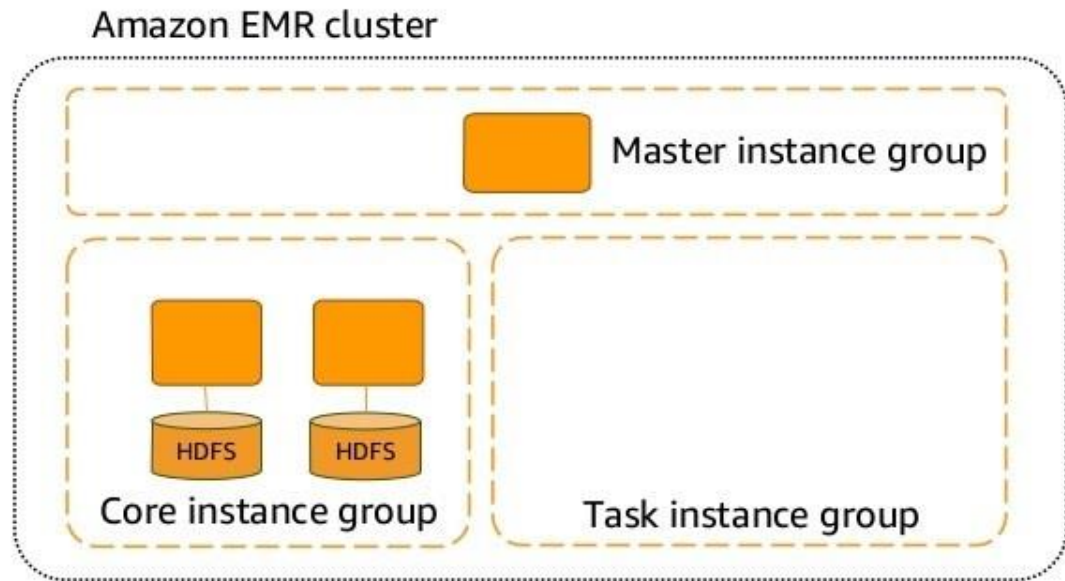
More memory

Jobs done faster



Amazon EMR Task nodes

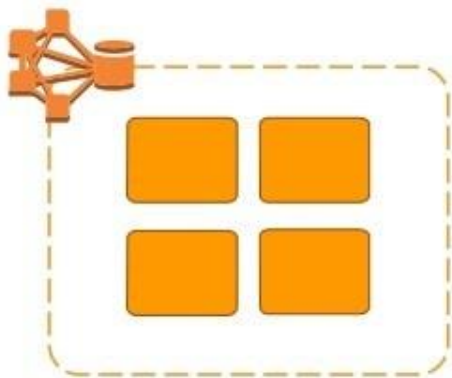
You can remove task nodes when processing is completed



Elastic clusters

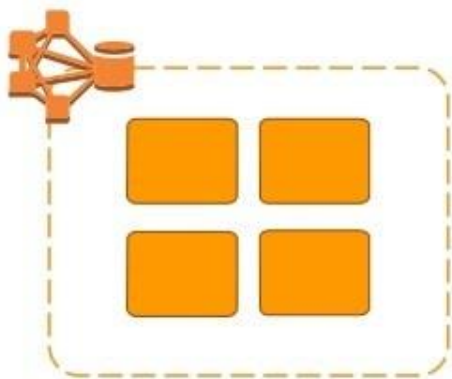
Elastic Clusters

1. Start cluster with certain number of nodes



Elastic Clusters

2. Set Auto-Scaling policies in CloudWatch



Elastic Clusters

3. Leverage EMR Auto-Scaling to add nodes or gracefully remove nodes based on CloudWatch policies



Scale up with Spot instances



10 node cluster running for 14 hours

Cost = \$1 * 10 * 14 = \$140

Resize nodes with Spot instances



Add 10 more nodes on Spot

Resize nodes with Spot instances



20 node cluster running for 7 hours

Cost = $\$1 * 10 * 7$ \$70

+ $\$0.5 * 10 * 7$ \$35

Total **\$105**

Resize nodes with Spot instances



50% less run-time: 14 → 7 hours

25% less cost: \$140 → \$105

Instance fleets for advanced Spot provisioning

Master Node



Core Instance Fleet



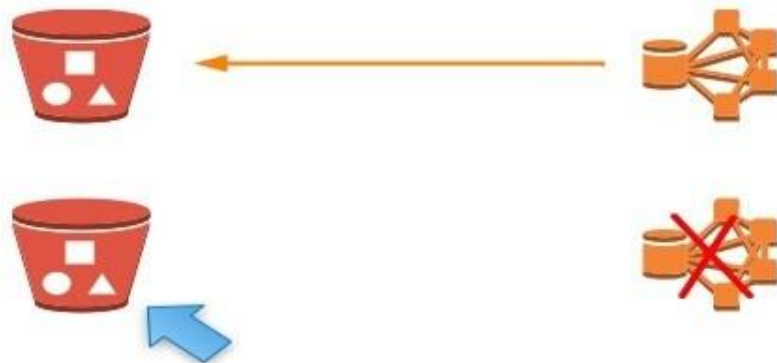
Task Instance Fleet



- Provision from a list of instance types with Spot and On-Demand
- Launch in the most optimal Availability Zone based on capacity/price
- Spot Block support (Defined Duration)

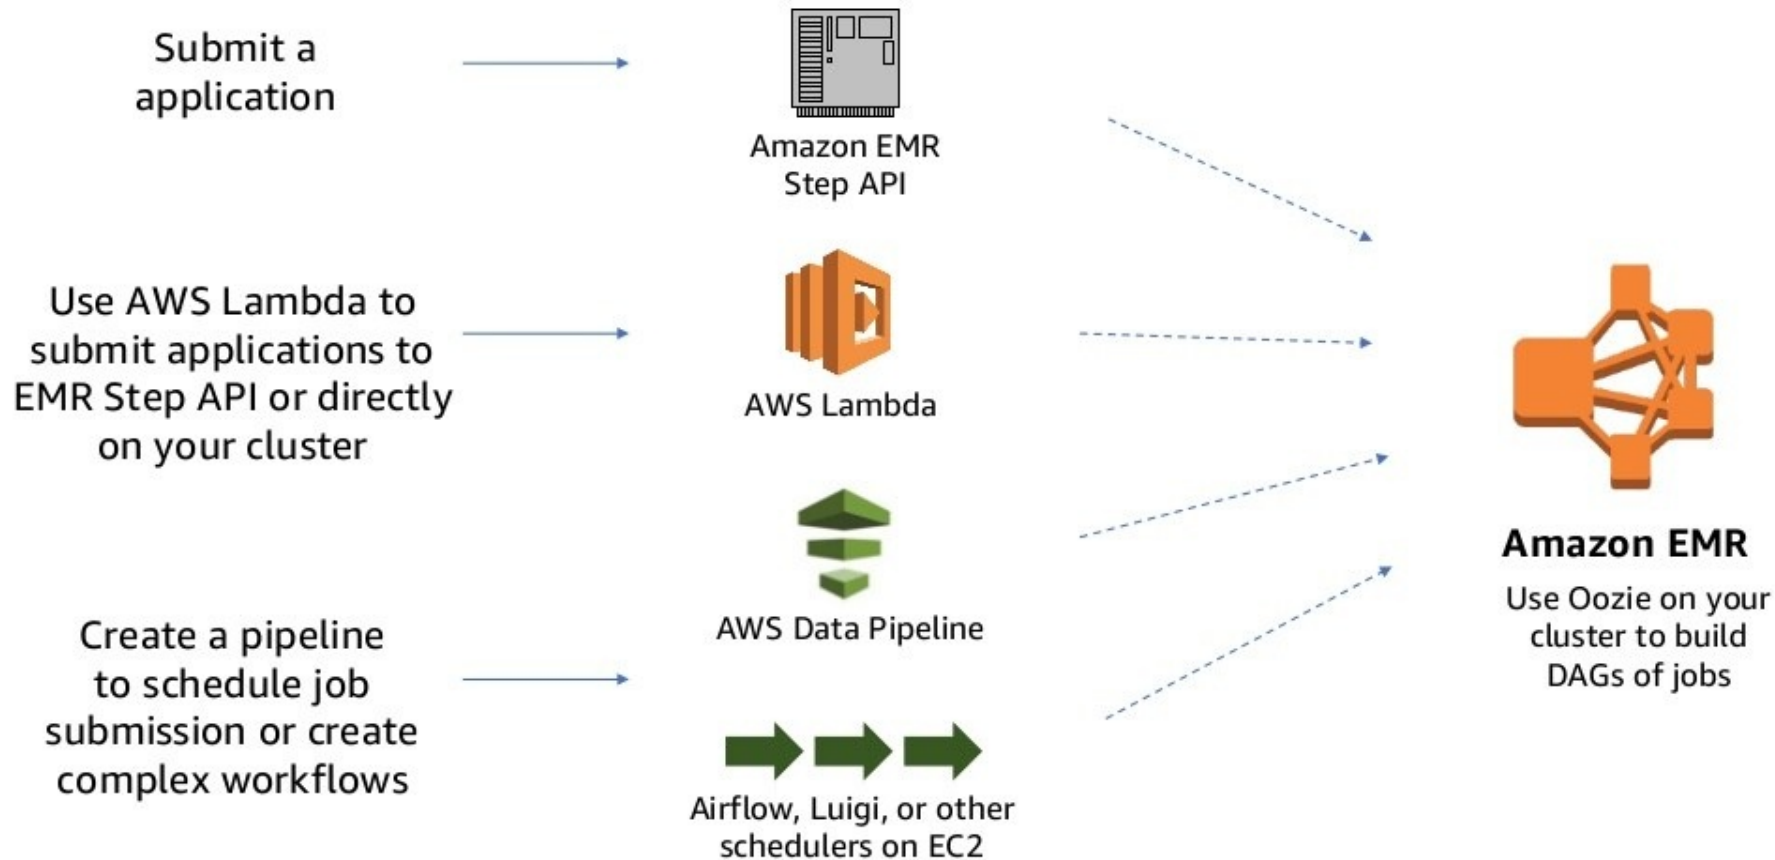
EMR Enables Transient Clusters

- Cluster lives for the **duration of the job**
- Shut down the cluster when the job is **done**
- Input and output data **persists** on Amazon S3



Data on Amazon S3

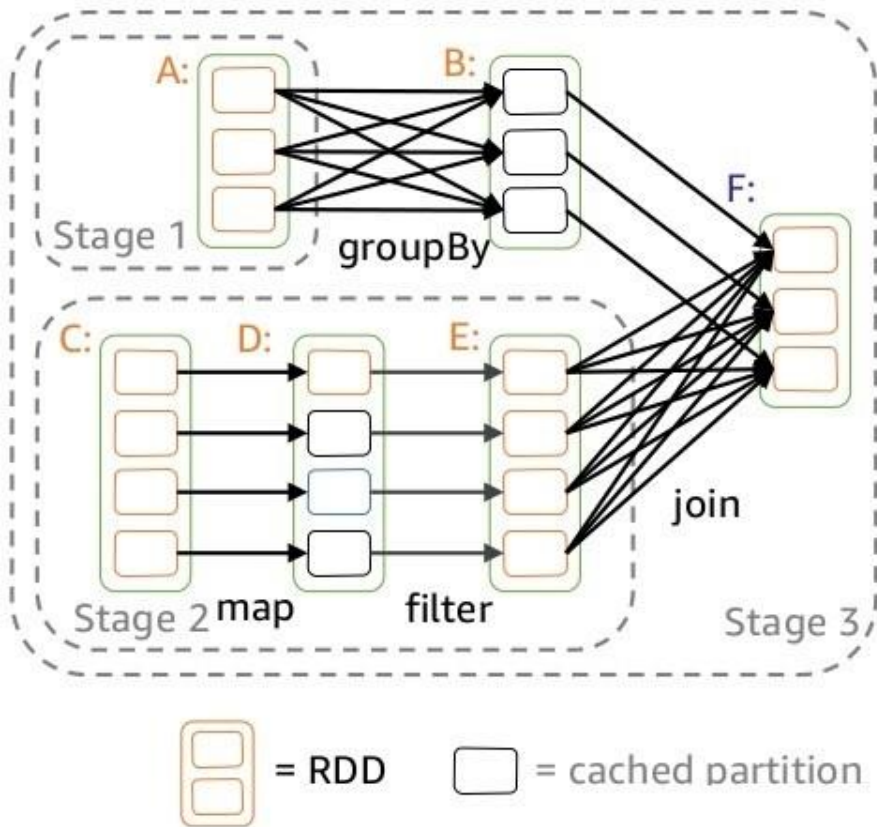
Options to submit jobs



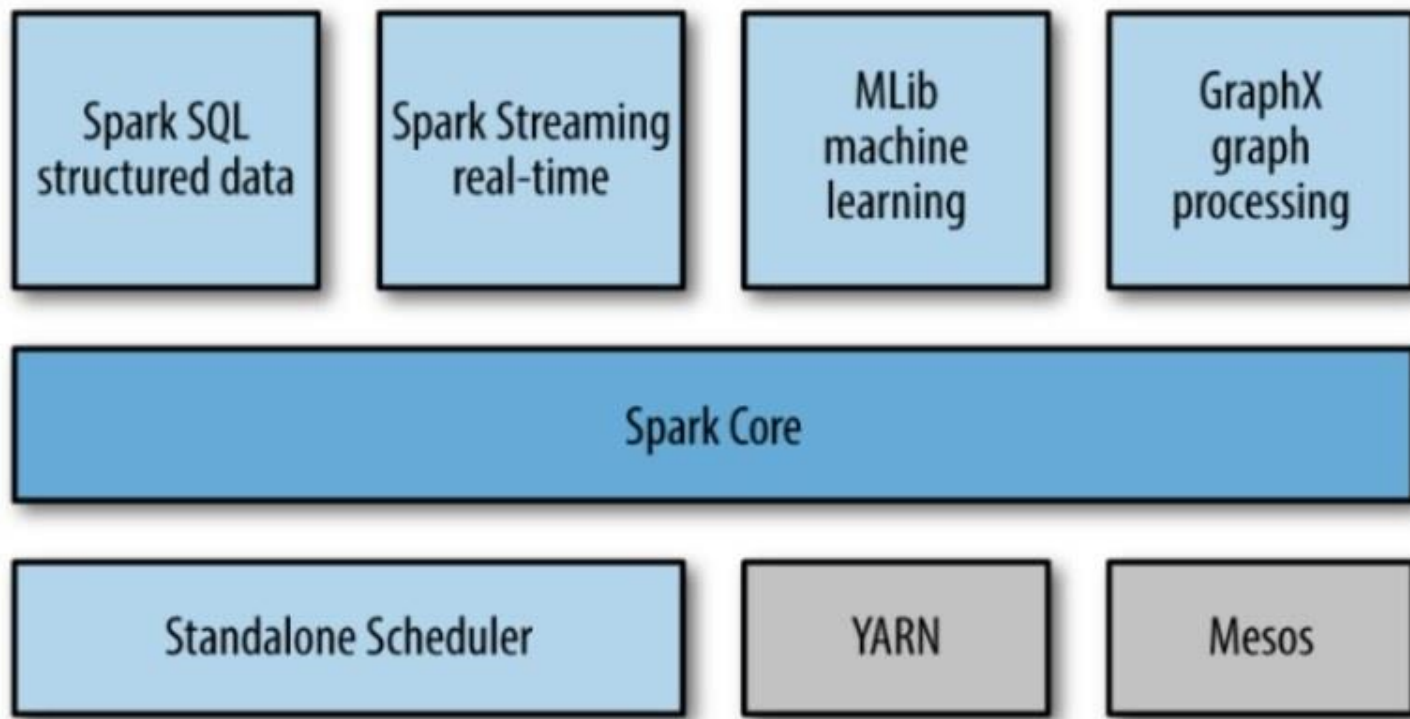
Beyond Map Reduce: Spark

Spark moves at interactive speed

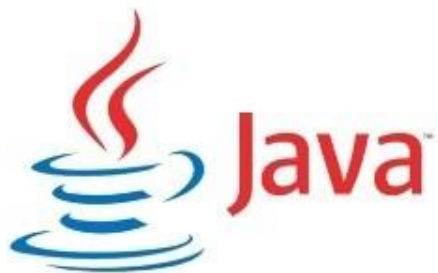
- Massively **parallel**
- Uses **DAGs** instead of map-reduce for execution
- Minimizes I/O by storing data in **memory**
- Partitioning-aware to avoid network-intensive shuffle



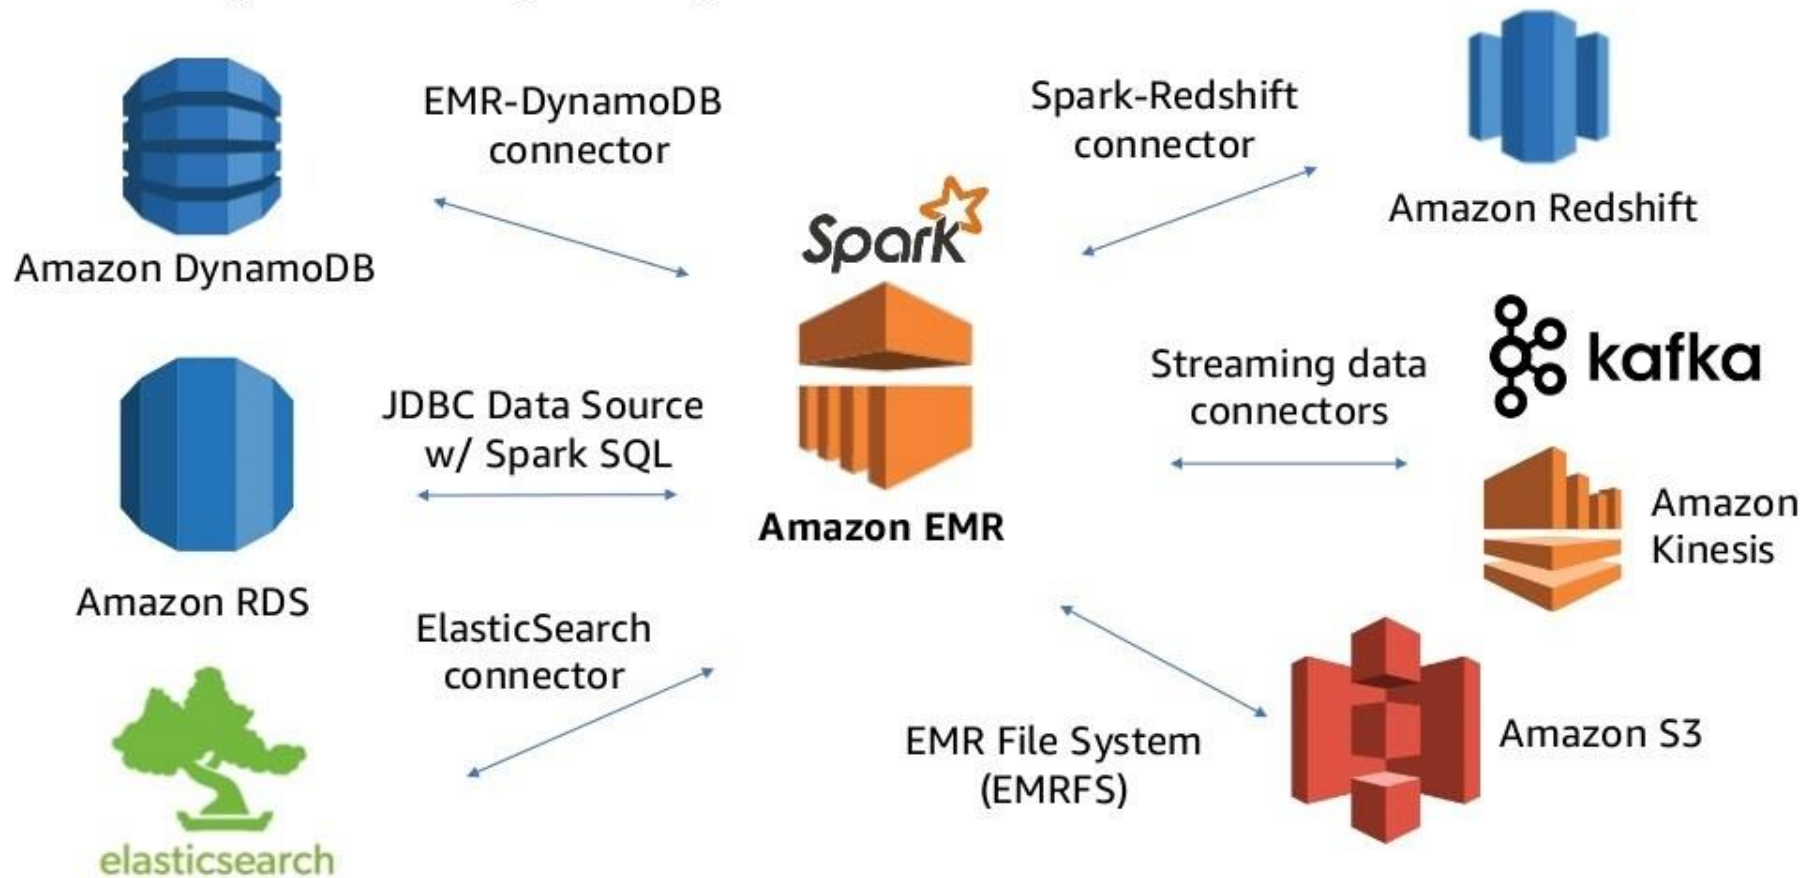
Spark components to match your use case



Spark speaks your language



Many storage layers to choose from



Use DataFrames for machine learning

```
// Prepare training documents from a list of (id, text, label) tuples
val training = sqlContext.createDataFrame(Seq(
  (0L, "a b c d e spark", 1.0),
  (1L, "b d", 0.0),
  (2L, "spark f g h", 1.0),
  (3L, "hadoop mapreduce", 0.0)
)).toDF("id", "text", "label")

// Configure an ML pipeline, which consists of three stages: tokenizer, hashingTF, and logisticRegression
val tokenizer = new Tokenizer()
  .setInputCol("text")
  .setOutputCol("words")
val hashingTF = new HashingTF()
  .setNumFeatures(1000)
  .setInputCol(tokenizer.getOutputCol)
  .setOutputCol("features")
val lr = new LogisticRegression()
  .setMaxIter(10)
  .setRegParam(0.01)
val pipeline = new Pipeline()
  .setStages(Array(tokenizer, hashingTF, lr))

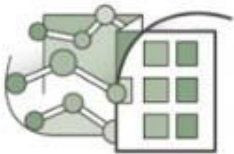
// Fit the pipeline to the training data
val model = pipeline.fit(training)
```

- Spark ML (replacing MLlib) uses DataFrames as input/output for models
- Create ML pipelines with a variety of distributed algorithms

Create DataFrames on streaming data



Amazon EMR



Easy to Use

Launch a cluster in minutes



Low Cost

Pay an hourly rate



Elastic

Easily add or remove capacity



Reliable

Spend less time
monitoring



Secure

Manage firewalls



Flexible

Customize the cluster

Resources

<https://aws.amazon.com/emr/>

<https://aws.amazon.com/blogs/bigdata>

<https://aws.amazon.com/whitepapers/>

<https://medium.com/@julsimon>



Thank you!

Data Pipeline

- A web service for scheduling regular data movement and data processing activities in the AWS cloud.
- Data Pipeline integrates with on-premise and cloud-based storage systems.
- A managed ETL (Extract-Transform-Load) service.
- Native integration with S3, DynamoDB, RDS, EMR, EC2 and Redshift.

Features

- You can quickly and easily provision pipelines that remove the development and maintenance effort required to manage your daily data operations, letting you focus on generating insights from that data.
- Data Pipeline provides built-in activities for common actions such as copying data between Amazon Amazon S3 and Amazon RDS, or running a query against Amazon S3 log data.
- Data Pipeline supports JDBC, RDS and Redshift databases.

Components

- A **pipeline definition** specifies the business logic of your data management.
- A **pipeline** schedules and runs tasks by creating EC2 instances to perform the defined work activities.
- Task Runner** polls for tasks and then performs those tasks. For example, Task Runner could copy log files to S3 and launch EMR clusters. Task Runner is installed and runs automatically on resources created by your pipeline definitions. You can write a custom task runner application, or you can use the Task Runner application that is provided by Data Pipeline.

Pipeline Definition

- From your pipeline definition, Data Pipeline determines the tasks, schedules them, and assigns them to task runners.
- If a task is not completed successfully, Data Pipeline retries the task according to your instructions and, if necessary, reassigns it to another task runner.
- If the task fails repeatedly, you can configure the pipeline to notify you.

Pipeline Definition

- A pipeline definition can contain the following types of components
- **Data Nodes** – The location of input data for a task or the location where output data is to be stored.
- **Activities** – A definition of work to perform on a schedule using a computational resource and typically input and output data nodes.
- **Preconditions** – A conditional statement that must be true before an action can run.

Pipeline Definition

And,

Scheduling Pipelines, Resources, Actions, Schedules

Task Runners

- When Task Runner is installed and configured, it polls Data Pipeline for tasks associated with pipelines that you have activated.
- When a task is assigned to Task Runner, it performs that task and reports its status back to Data Pipeline.